

Azat Dovgeldiyev

CSC 555

Assignment 5

Problem 1.

- a) Highly structured multi-table data that requires enforcing data constraints.
Relational database.
- b) Stock market data ticker with decisions that must be made in real time.
Streaming engine.
- c) LinkedIn type data with interconnected nodes where much of the information resides in the links between nodes.
Graph database.
- d) An image storage system that allows lookup images by file name.
Key-Value stores
- e) A collection of JSON objects (e.g., tweets).
Document-oriented store.
- f) Data that is stored in large sparse tables that are continuously growing (new rows/columns).
Column-oriented store

Problem 2.

- a) Compute the page rank for the nodes in this graph. If you are multiplying matrices manually, you may stop computing after 5 steps. If you use a tool (e.g., Matlab, python) for matrix multiplication, you should get your answer to converge.

M =

0	1	1/3	1	A
0	0	1/3	0	B
1	0	0	0	Y
0	0	1/3	0	Z
A	B	Y	Z	

v =

1/4
1/4
1/4
1/4

5 steps, $M^5 \cdot v \Rightarrow$

Step 1:

$$\begin{aligned}\left(0 * \frac{1}{4}\right) + \left(1 * \frac{1}{4}\right) + \left(\frac{1}{3} * \frac{1}{4}\right) + \left(1 * \frac{1}{4}\right) &= \frac{7}{12} \\ \left(0 * \frac{1}{4}\right) + \left(0 * \frac{1}{4}\right) + \left(\frac{1}{3} * \frac{1}{4}\right) + \left(0 * \frac{1}{4}\right) &= \frac{1}{12} \\ \left(1 * \frac{1}{4}\right) + \left(0 * \frac{1}{4}\right) + \left(0 * \frac{1}{4}\right) + \left(0 * \frac{1}{4}\right) &= \frac{1}{4} \\ \left(0 * \frac{1}{4}\right) + \left(0 * \frac{1}{4}\right) + \left(\frac{1}{3} * \frac{1}{4}\right) + \left(0 * \frac{1}{4}\right) &= \frac{1}{12}\end{aligned}$$

Step 2:

$$\begin{aligned}\left(0 * \frac{7}{12}\right) + \left(1 * \frac{1}{12}\right) + \left(\frac{1}{3} * \frac{1}{4}\right) + \left(1 * \frac{1}{12}\right) &= \frac{1}{4} \\ \left(0 * \frac{7}{12}\right) + \left(0 * \frac{1}{12}\right) + \left(\frac{1}{3} * \frac{1}{4}\right) + \left(0 * \frac{1}{12}\right) &= \frac{1}{12} \\ \left(1 * \frac{7}{12}\right) + \left(0 * \frac{1}{12}\right) + \left(0 * \frac{1}{4}\right) + \left(0 * \frac{1}{12}\right) &= \frac{7}{12} \\ \left(0 * \frac{7}{12}\right) + \left(0 * \frac{1}{12}\right) + \left(\frac{1}{3} * \frac{1}{4}\right) + \left(0 * \frac{1}{12}\right) &= \frac{1}{12}\end{aligned}$$

Step 3:

$$\begin{aligned}\left(0 * \frac{1}{4}\right) + \left(1 * \frac{1}{12}\right) + \left(\frac{1}{3} * \frac{7}{12}\right) + \left(1 * \frac{1}{12}\right) &= \frac{13}{36} \\ \left(0 * \frac{1}{4}\right) + \left(0 * \frac{1}{12}\right) + \left(\frac{1}{3} * \frac{7}{12}\right) + \left(0 * \frac{1}{12}\right) &= \frac{7}{36} \\ \left(1 * \frac{1}{4}\right) + \left(0 * \frac{1}{12}\right) + \left(0 * \frac{7}{12}\right) + \left(0 * \frac{1}{12}\right) &= \frac{1}{4} \\ \left(0 * \frac{1}{4}\right) + \left(1 * \frac{1}{12}\right) + \left(\frac{1}{3} * \frac{7}{12}\right) + \left(1 * \frac{1}{12}\right) &= \frac{7}{36}\end{aligned}$$

Step 4:

$$\begin{aligned}\left(0 * \frac{13}{36}\right) + \left(1 * \frac{7}{36}\right) + \left(\frac{1}{3} * \frac{1}{4}\right) + \left(1 * \frac{7}{36}\right) &= \frac{17}{36} \\ \left(0 * \frac{13}{36}\right) + \left(0 * \frac{7}{36}\right) + \left(\frac{1}{3} * \frac{1}{4}\right) + \left(0 * \frac{7}{36}\right) &= \frac{1}{12} \\ \left(1 * \frac{13}{36}\right) + \left(0 * \frac{7}{36}\right) + \left(0 * \frac{1}{4}\right) + \left(0 * \frac{7}{36}\right) &= \frac{13}{36} \\ \left(0 * \frac{13}{36}\right) + \left(0 * \frac{7}{36}\right) + \left(\frac{1}{3} * \frac{1}{4}\right) + \left(0 * \frac{7}{36}\right) &= \frac{1}{12}\end{aligned}$$

Step 5:

$$\begin{aligned}\left(0 * \frac{17}{36}\right) + \left(1 * \frac{1}{12}\right) + \left(\frac{1}{3} * \frac{13}{36}\right) + \left(1 * \frac{1}{12}\right) &= \frac{31}{108} \\ \left(0 * \frac{17}{36}\right) + \left(0 * \frac{1}{12}\right) + \left(\frac{1}{3} * \frac{13}{36}\right) + \left(0 * \frac{1}{12}\right) &= \frac{13}{108}\end{aligned}$$

$$\left(1 * \frac{17}{36}\right) + \left(0 * \frac{1}{12}\right) + \left(0 * \frac{13}{36}\right) + \left(0 * \frac{1}{12}\right) = \frac{17}{36}$$

$$\left(0 * \frac{17}{36}\right) + \left(0 * \frac{1}{12}\right) + \left(\frac{1}{3} * \frac{13}{36}\right) + \left(0 * \frac{1}{12}\right) = \frac{13}{108}$$

Rank :

31/108
13/108
17/36
13/108

b) Now consider a graph with dead-end nodes Q and P:

What is the page rank of Q?

Since p and q are the dead ends, we can drop them:

M =

V =

0	1	½	A
0	0	½	Y
1	0	0	Z
A	Y	Z	

1/3
1/3
1/3

After calculating 5 steps:

V =

3/8
5/24
5/12

Pagerank of Q =

$$\left(0 * \frac{3}{8}\right) + \left(\frac{5}{24} * \frac{1}{2}\right) + \left(\frac{5}{12} * \frac{1}{3}\right) = \frac{35}{144} = 0.243$$

What is the page rank of P?

Now, we can compute the PageRank for P. The node has only one predecessor, Q, and Q has only one successor. Thus, the PageRank of P is the same as that of Q.

$$1 * 35/144 = 0.243$$

c) If we eliminate all the dead ends, remaining root node has only link to itself and rank 1.

First dead end has only the first node as its predecessor, and root node has two successors, so the contribution will be 1/2 to the first dead end, and rank for the first node is 1/2 (as $1 * 1/2$), therefore subsequent dead-end nodes have ranks 1/2.

Problem 3.

- a) What will the Hive query “compute average price” return? (yes, this question is as obvious as it seems, asked for comparison with part-b and part-c)

```
data = [('1pm', 6), ('2pm', 15), ('3pm', 15),  
        ('4pm', 20), ('5pm', 10), ('6pm', 20),  
        ('7pm', 20), ('8pm', 24), ('9pm', 23),  
        ('10pm', 32), ('11pm', 26), ('12am', 40)]
```

```
#a hive average price  
price = 0  
for k,v in data:  
    price+=v  
print("average price: ",price/len(data))
```

average price: 20.916666666666668

- b) What will a Storm streaming query “compute average price per each 3 hour window” return? (tumbling, i.e., non-overlapping window of tuples). For example, the first window would 1pm-4pm. Second window would be 4pm-7pm. If you are wondering about overlap, I would recommend defaulting to (1pm-4pm] (4pm-7pm].

```
# 3 hour window  
price = 0  
lst = []  
c = 1  
for k, v in data:  
    if c%3 == 0:  
        price+=v  
        lst.append(k)  
        print(lst,price/3)  
        c=1  
        price = 0  
        lst=[]  
    else:  
        c+=1  
        price+=v  
        lst.append(k)
```

```
['1pm', '2pm', '3pm'] 12.0  
['4pm', '5pm', '6pm'] 16.666666666666668  
['7pm', '8pm', '9pm'] 22.333333333333332  
['10pm', '11pm', '12am'] 32.666666666666664
```

- c) What will a Storm query “compute average price per each 3 hour window” return? (sliding, i.e. overlapping window of tuples, moving the window forward 2 hours each time). First window is 1pm-4pm, second window is 3pm-6pm

```
# 3 hour window -sliding
price = 0
lst = []
c = 1
for k, v in data:
    if c%3 == 0:
        price+=v
        lst.append(k)
        print(lst,price/3)
        c=2
        price = 0
        lst=[k]
    else:
        c+=1
        price+=v
        lst.append(k)
```

```
['1pm', '2pm', '3pm'] 12.0
['3pm', '4pm', '5pm'] 10.0
['5pm', '6pm', '7pm'] 13.333333333333334
['7pm', '8pm', '9pm'] 15.666666666666666
['9pm', '10pm', '11pm'] 19.333333333333332
```

Problem 4.

Run another custom MapReduce job, implementing a solution for the following query:
For Employee(EID, EFirst, ELast, Phone) and Customer(CID, CFirst, CLast, Address), find everyone with the same name using MapReduce:
find . -name “Hadoop-streaming-2.6.4.jar” -print
cp .

```
SELECT EFirst, ELast, Phone, Address
FROM Employee, Customer
WHERE EFirst = CFirst AND ELast = CLast
```

eMapper.py

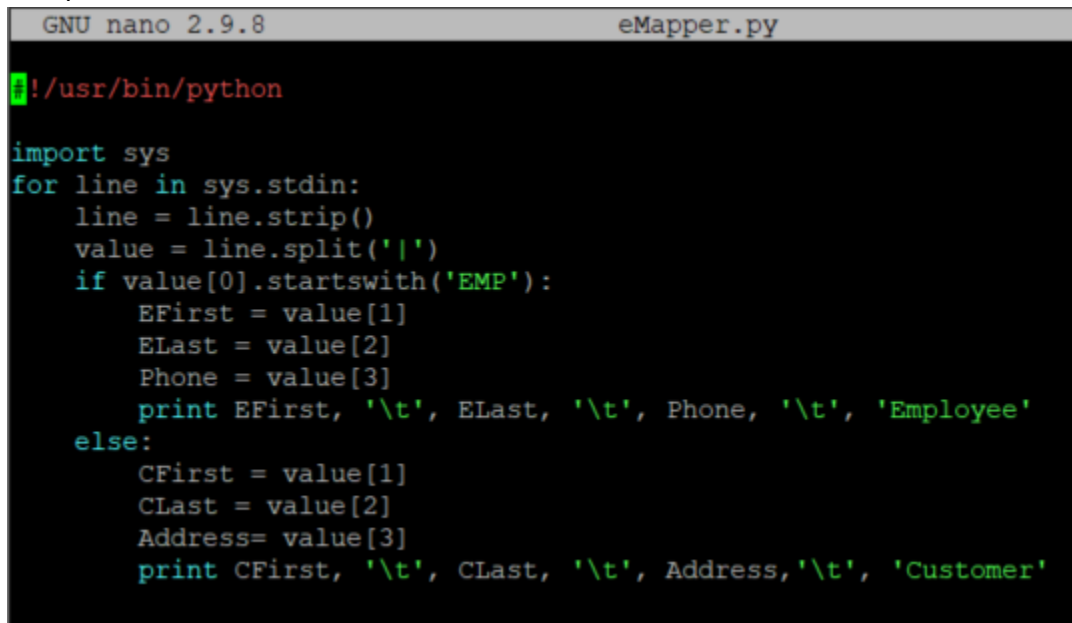
```
#!/usr/bin/python
```

```
import sys
for line in sys.stdin:
    line = line.strip()
    value = line.split('|')
    if value[0].startswith('EMP'):
        EFirst = value[1]
        ELast = value[2]
```

```

    Phone = value[3]
    print EFirst, '\t', ELast, '\t', Phone, '\t', 'Employee'
else:
    CFirst = value[1]
    CLast = value[2]
    Address= value[3]
    print CFirst, '\t', CLast, '\t', Address, '\t', 'Customer'

```



```

GNU nano 2.9.8 eMapper.py
#!/usr/bin/python

import sys
for line in sys.stdin:
    line = line.strip()
    value = line.split('|')
    if value[0].startswith('EMP'):
        EFirst = value[1]
        ELast = value[2]
        Phone = value[3]
        print EFirst, '\t', ELast, '\t', Phone, '\t', 'Employee'
    else:
        CFirst = value[1]
        CLast = value[2]
        Address= value[3]
        print CFirst, '\t', CLast, '\t', Address, '\t', 'Customer'

```

eReducer.py

```
#!/usr/bin/python
```

```
import sys
```

```
cKey = None
```

```
EFirst = None
```

```
ELast = None
```

```
Phone = None
```

```
Address = None
```

```
for line in sys.stdin:
```

```
    line = line.strip()
```

```
    value = line.split('\t')
```

```
    k = value[0] + ' ' + value[1]
```

```
    v = '\t'.join(value[2:])
```

```
    if cKey == k:
```

```
        if v.endswith('Employee'):
```

```

        EFirst = value[0]
        ELast = value[1]
        Phone = value[2]
        if v.endswith('Customer'):
            Address = value[2]
    else:
        if cKey:
            lp = len(Phone)
            la = len(Address)
            if (lp * la > 0):
                print EFirst, '\t', ELast, '\t', Phone, '\t', Address
        EFirst = ""
        ELast = ""
        Phone = ""
        Address = ""
        cKey = k
        if v.endswith('Employee'):
            EFirst = value[0]
            ELast = value[1]
            Phone = value[2]
            Address = ""
        elif v.endswith('Customer'):
            EFirst = ""
            ELast = ""
            Phone = ""
            Address = value[2]
            lp = len(Phone)
            la = len(Address)
            if (lp * la > 0):
                print EFirst, '\t', ELast, '\t', Phone, '\t', Address
    lp = len(Phone)
    la = len(Address)
    if (lp * la > 0):
        print EFirst, '\t', ELast, '\t', Phone, '\t', Address

```

```
hadoop jar hadoop-streaming-2.6.4.jar-D stream.num.map.output.key.fields=2 -input
/user/ec2-user/ecust -output /data/ant5 -mapper eMapper.py -reducer eReducer.py -file
eMapper.py -file eReducer.py
```

```
Map input records=110000
Map output records=110000
Map output bytes=6848671
Map output materialized bytes=7068689
Input split bytes=315
Combine input records=0
Combine output records=0
Reduce input groups=14225
Reduce shuffle bytes=7068689
Reduce input records=110000
Reduce output records=188
Spilled Records=220000
Shuffled Maps =3
Failed Shuffles=0
Merged Map outputs=3
GC time elapsed (ms)=429
CPU time spent (ms)=4200
Physical memory (bytes) snapshot=942555136
Virtual memory (bytes) snapshot=8511180800
Total committed heap usage (bytes)=678952960

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
Bytes Read=6311657
File Output Format Counters
Bytes Written=12101
20/11/08 23:32:36 INFO streaming.StreamJob: Output directory: /data/ant5
```

```
[ec2-user@ip-172-31-74-226 hadoop-2.6.4]$ hadoop fs -ls /data/ant5
Found 2 items
-rw-r--r--  2 ec2-user supergroup          0 2020-11-08 23:32 /data/ant5/_SUCCESS
-rw-r--r--  2 ec2-user supergroup    12101 2020-11-08 23:32 /data/ant5/part-000000
[ec2-user@ip-172-31-74-226 hadoop-2.6.4]$ hadoop fs -cat /data/ant5/part-000000
Brendan      Anastasio      70      613 Devon Court, West Orange, NJ 07052
Brendan      Berenbaum      76      343 Franklin Street, Fort Walton Beach,
FL 32547
Brendan      Bosque      79      783 8th Avenue, Elkton, MD 21921
Brendan      Cashin      71      742 Beechwood Drive, Fairfax, VA 22030
Brendan      Lembke      32      926 Olive Street, Fort Wayne, IN 46804
Brendan      Mabe      41      761 Route 5, Chandler, AZ 85224
Brendan      Maynor      80      693 Orchard Street, Algonquin, IL 60102
Brendan      Mcdougale    24      228 Homestead Drive, Aiken, SC 29803
Brendan      Mullican     40      992 Oxford Court, Tewksbury, MA 01876
Brendan      Platt      45      232 Old York Road, Englewood, NJ 07631
Brendan      Read      66      222 Sycamore Lane, Garden City, NY 11530
Brendan      Tyrrell     50      455 Warren Street, Wyandotte, MI 48192
Brendan      Walpole     38      949 Maple Street, Oakland Gardens, NY 11
364
Francoise   Anastasio      73      52 Augusta Drive, Clayton, NC 27520
Francoise   Berenbaum      46      957 Liberty Street, Satellite Beach, FL
```


32937			
Francoise	Bosque	34	228 Dogwood Drive, Melrose, MA 02176
Francoise	Cashin	80	80 Dogwood Drive, Fairhope, AL 36532
Francoise	Hartley	81	687 Cedar Street, Elgin, IL 60120
Francoise	Lembke	72	55 Forest Avenue, Media, PA 19063
Francoise	Mabe	52	655 Park Avenue, Waynesboro, PA 17268
Francoise	Maynor	62	128 Summit Avenue, Cranston, RI 02920
Francoise	Mcdougale	80	761 Route 5, Chandler, AZ 85224
Francoise	Mullican	61	274 Meadow Street, El Paso, TX 79930
Francoise	Platt	34	803 Cedar Lane, Essex, MD 21221
Francoise	Read	33	589 Bridge Street, Fort Worth, TX 76110
Francoise	Tyrrell	39	969 Valley View Road, Deland, FL 32720
Francoise	Walpole	47	957 Liberty Street, Satellite Beach, FL
32937			
Freeda	Anastasio	69	480 Strawberry Lane, South Lyon, MI 4817
8			
Freeda	Berenbaum	48	176 Warren Street, Piqua, OH 45356
Freeda	Bosque	70	687 Cedar Street, Elgin, IL 60120
Freeda	Cashin	73	1 Homestead Drive, Willoughby, OH 44094
Freeda	Hartley	32	187 Hickory Lane, Raleigh, NC 27603
Freeda	Lembke	36	228 Dogwood Drive, Melrose, MA 02176
Freeda	Mabe	37	797 Cedar Street, Muskegon, MI 49441
Freeda	Maynor	80	516 Essex Court, Adrian, MI 49221
Freeda	Mcdougale	58	783 8th Avenue, Elkton, MD 21921
Freeda	Mullican	50	517 Andover Court, Naugatuck, CT 06770
Freeda	Platt	24	187 Hickory Lane, Raleigh, NC 27603
Freeda	Read	66	949 Maple Street, Oakland Gardens, NY 11364
Freeda	Tyrrell	29	480 Strawberry Lane, South Lyon, MI 4817
8			
Freeda	Walpole	24	688 Main Street West, Alexandria, VA 223
04			
Hosea	Anastasio	68	295 Hillcrest Drive, Green Bay, WI 54302
Hosea	Berenbaum	23	720 Route 20, Los Banos, CA 93635
Hosea	Bosque	64	187 Magnolia Avenue, Maryville, TN 37803
Hosea	Cashin	65	957 Liberty Street, Satellite Beach, FL 32937
Hosea	Hartley	25	635 Cross Street, Monsey, NY 10952
Hosea	Lembke	34	655 Park Avenue, Waynesboro, PA 17268
Hosea	Mabe	47	274 Meadow Street, El Paso, TX 79930
Hosea	Maynor	58	705 Main Street South, Anaheim, CA 92806
Hosea	Mcdougale	28	187 Magnolia Avenue, Maryville, TN 37803
Hosea	Mullican	75	295 Hillcrest Drive, Green Bay, WI 54302
Hosea	Platt	23	455 Warren Street, Wyandotte, MI 48192
Hosea	Read	55	455 Warren Street, Wyandotte, MI 48192
Hosea	Tyrrell	71	800 Rosewood Drive, Soddy Daisy, TN 37379
Hosea	Walpole	28	477 Pine Street, Neenah, WI 54956
Isidro	Anastasio	25	55 Forest Avenue, Media, PA 19063
Isidro	Berenbaum	24	687 Cedar Street, Elgin, IL 60120
Isidro	Bosque	63	128 Summit Avenue, Cranston, RI 02920
Isidro	Cashin	26	705 Main Street South, Anaheim, CA 92806
Isidro	Hartley	38	253 Route 2, Sun Prairie, WI 53590
Isidro	Mabe	71	187 Magnolia Avenue, Maryville, TN 37803
Isidro	Maynor	79	538 6th Street West, Springboro, OH 4506
6			
Isidro	Mcdougale	76	844 Marshall Street, Oakland Gardens, NY
11364			
Isidro	Mullican	25	2 Pennsylvania Avenue, Winchester, VA 22
601			
Isidro	Platt	26	253 Route 2, Sun Prairie, WI 53590
Isidro	Read	25	103 Oxford Court, Newark, NJ 07103

Problem 5.

a)

```
[ec2-user@ip-172-31-74-226 ~]$ less web-Stanford.txt
# Directed graph (each unordered pair of nodes is saved once): web-Stanford.txt
# Stanford web graph from 2002
# Nodes: 281903 Edges: 2312497
# FromNodeId    ToNodeId
```

There are 281903 nodes and 2312497 edges the file contains.

b) 4m20.417s

```
total committed heap usage (bytes): 4800000
Shuffle Errors
      BAD_ID=0
      CONNECTION=0
      IO_ERROR=0
      WRONG_LENGTH=0
      WRONG_MAP=0
      WRONG_REDUCE=0
File Input Format Counters
      Bytes Read=22629726
File Output Format Counters
      Bytes Written=7333875
DONE!

real    4m20.417s
user    0m6.774s
sys     0m0.493s
[ec2-user@ip-172-31-74-226 src]$
```

c)

```
[ec2-user@ip-172-31-74-226 src]$ hadoop fs -cat /data/prOutput/result/part-r-00000 | more
0.15000000596046448      75380
0.15000000596046448      155237
0.15000000596046448      75378
0.15000000596046448      155226
0.15000000596046448      155222
0.15000000596046448      155221
0.15000000596046448      75372
0.15000000596046448      125860
0.15000000596046448      47294
0.15000000596046448      155216
0.15000000596046448      155204
0.15000000596046448      155203
0.15000000596046448      47303
0.15000000596046448      155116
0.15000000596046448      125881
0.15000000596046448      47316
0.15000000596046448      75340
0.15000000596046448      47317
0.15000000596046448      125872
0.15000000596046448      75358
0.15000000596046448      75357
0.15000000596046448      47327
0.15000000596046448      47336
--More--
```