

Node	Last contact	Under replicated blocks	Blocks with no live replicas	Under Replicated Blocks In files under construction
------	--------------	-------------------------	------------------------------	--

Security is off.

Safemode is off.

1 files and directories, 0 blocks = 1 total filesystem object(s).

Heap Memory used 136.12 MB of 195 MB Heap Memory. Max Heap Memory is 889 MB.

Non Heap Memory used 39.61 MB of 40.31 MB Committed Non Heap Memory. Max Non Heap Memory is -1 B.

Configured Capacity:	89.96 GB
DFS Used:	12 KB
Non DFS Used:	5.79 GB
DFS Remaining:	84.17 GB
DFS Used%:	0%
DFS Remaining%:	93.56%
Block Pool Used:	12 KB
Block Pool Used%:	0%
DataNodes usages% (Min/Median/Max/stdDev):	0.00% / 0.00% / 0.00% / 0.00%
Live Nodes	3 (Decommissioned: 0)
Dead Nodes	0 (Decommissioned: 0)
Decommissioning Nodes	0

Repeat the steps for wordcount using bioproject.xml from Assignment 1 and submit screenshots of running it.

```
Shuffle Errors
      BAD_ID=0
      CONNECTION=0
      IO_ERROR=0
      WRONG_LENGTH=0
      WRONG_MAP=0
      WRONG_REDUCE=0
File Input Format Counters
      Bytes Read=231153099
File Output Format Counters
      Bytes Written=20056175

real    0m42.283s
user    0m3.791s
sys     0m0.232s
[ec2-user@ip-172-31-74-226 ~]$
```

Size of wordcount

```
[ec2-user@ip-172-31-74-226 ~]$ hadoop fs -du /data/wordcount1/
0          /data/wordcount1/_SUCCESS
20056175   /data/wordcount1/part-r-00000
[ec2-user@ip-172-31-74-226 ~]$
```

```

arctica</Name>      5
arctica</OrganismName> 5
arcticus            31
arcticus<lt;/i>gt;    2
arcticus</Name>      4
arcticus</OrganismName> 4
holarctica          77
humans.Antarctic     1
palearctica          66
palearctica</Name>    1
sub-Antarctic        4
sub-arctic           4
subantarctic         1
subantarcticus       7
subantarcticus</Name> 1
subantarcticus</OrganismName> 1
subarctic            21
[ec2-user@ip-172-31-74-226 ~]$

```

Comparing to the first assignment where time completed was about 1 minute and 14 seconds, this time it is faster with time completed for 42 seconds, and the reason for that one node is slower than working with 3 nodes.

## Part 2: Hive

- 1) Run the following query in Hive and report the time it takes to execute:

```

19961225      594517529
19961228      555549994
19961231      612180222
Time taken: 39.587 seconds, Fetched: 366 row(s)
hive>

```

- 2) Perform the following transform operation using SELECT TRANSFORM on the dwdate table by creating a new table. The new dwdate table will combine d\_daynuminweek, d\_daynuminmonth, and d\_daynuminyear into a single column in the new table. You should also eliminate of the last 4 columns (d\_lastdayinweekfl, d\_lastdayinmonthfl, d\_holidayfl, and d\_weekdayfl). The final table will have 6 fewer columns than the original table because you merge 3 columns into 1 and remove 4 columns.

I created a new table , last column d\_weekmonth will combine 3 columns:

```

hive> create table new_dwdate (
  >   d_datekey          int,
  >   d_date             varchar(19),
  >   d_dayofweek        varchar(10),
  >   d_month            varchar(10),
  >   d_year             int,
  >   d_yearmonthnum     int,
  >   d_yearmonth        varchar(8),
  >   d_monthnuminyear   int,
  >   d_weeknuminyear    int,
  >   d_sellingseason    varchar(13),
  >   d_weekmonth        varchar(8)
  > )
  > ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' STORED AS TEXTFILE;
OK
Time taken: 0.078 seconds

```

```
INSERT OVERWRITE TABLE new_dwdate SELECT TRANSFORM (d_datekey, d_date, d_dayofweek,
d_month, d_year, d_yearmonthnum,d_yearmonth,d_daynuminweek, d_daynuminmonth,
d_daynuminyear, d_monthnuminyear, d_weeknuminyear, d_sellingseason, d_lastdayinweekfl,
d_lastdayinmonthfl, d_holidayfl, d_weekdayfl) USING 'python pdate1.py' AS (d_datekey, d_date,
d_dayofweek, d_month, d_year, d_yearmonthnum,d_yearmonth, d_monthnuminyear,
d_weeknuminyear, d_sellingseason, d_weekmonth) FROM dwdate;
```

## TESTING;

```
Moving data to: hdfs://172.31.74.226/user/hive/warehouse/new_dwdate/.hive-staging_hive_2020-10-25_21-23-48_023_
Loading data to table default.new_dwdate
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 2.08 sec HDFS Read: 239794 HDFS Write: 207041 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 80 msec
OK
Time taken: 12.254 seconds
hive> select * from new_dwdate limit 5;
OK
19920101 January 1, 1992 Thursday January 1992 199201 Jan1992 1 1 Winter 5_1_1
19920102 January 2, 1992 Friday January 1992 199201 Jan1992 1 1 Winter 6_2_2
19920103 January 3, 1992 Saturday January 1992 199201 Jan1992 1 1 Winter 7_3_3
19920104 January 4, 1992 Sunday January 1992 199201 Jan1992 1 1 Winter 1_4_4
19920105 January 5, 1992 Monday January 1992 199201 Jan1992 1 1 Winter 2_5_5
Time taken: 0.067 seconds, Fetched: 5 row(s)
hive> select * from new_dwdate limit 5;
```

```
[ec2-user@ip-172-31-74-226 ~]$ hadoop fs -ls /user/hive/warehouse/new_dwdate
Found 1 items
-rwxrwxr-x 2 ec2-user supergroup 206961 2020-10-25 21:23 /user/hive/warehouse/new_dwdate/000000_0
```

Python code :

```
GNU nano 2.9.8 pdate1.py

#!/usr/bin/python
import sys

#read
for line in sys.stdin:
    line=line.strip()
    vals=line.split('\t')
    week=vals[7]
    month=vals[8]
    year=vals[9]
    vals.append(week+'_'+month+'_'+year)
    print('\t'.join([vals[0],vals[1],vals[2],vals[3],vals[4],vals[5],vals[6],vals[10],vals[11],vals[12],vals[17]]))
```

```
#!/usr/bin/python
import sys

#read
for line in sys.stdin:
    line=line.strip()
    vals=line.split('\t')
    week=vals[7]
    month=vals[8]
    year=vals[9]
    vals.append(week+'_'+month+'_'+year)

print('\t'.join([vals[0],vals[1],vals[2],vals[3],vals[4],vals[5],vals[6],vals[10],vals[11],vals[12],vals[17]]))
```

## Part 3: Pig

```
SELECT lo_discount, COUNT(lo_extendedprice)
FROM lineorder
GROUP BY lo_discount;
```

```
lod = LOAD '/user/ec2-user/lineorder.tbl' USING PigStorage('|') AS(lo_orderkey:int,lo_linenummer:int,
lo_custkey:int, lo_partkey:int, lo_suppkey:int, lo_orderdate:int, lo_orderpriority:chararray,
lo_shippriority:chararray, lo_quantity:int, lo_extendedprice:int, lo_ordertotalprice:int, lo_discount:int,
lo_revenue:int,lo_supplycost:int, lo_tax:int, lo_commitdate:int, lo_shipmode:chararray);
```

```
by_discount = group lod by lo_discount;
filter_discount = FOREACH by_discount GENERATE group as
lo_discount,COUNT(lod.lo_extendedprice);
dump filter_discount;
```

```
Success!

Job Stats (time in seconds):
JobId  Maps   Reduces MaxMapTime   MinMapTime   AvgMapTime   MedianMapTime   MaxReduceTime   MinReduceTime   AvgReduceTime   M
cetime Alias   Feature Outputs
job_1603645671591_0013_5  1  52  29  47  51  20  20  20  20  by_discount,filter_discount,lod G
OMBINER hdfs://172.31.74.226/tmp/temp-956455315/tmp-967075946,

Input(s):
Successfully read 6001215 records (594331260 bytes) from: "/user/ec2-user/lineorder.tbl"

Output(s):
Successfully stored 11 records (119 bytes) in: "hdfs://172.31.74.226/tmp/temp-956455315/tmp-967075946"

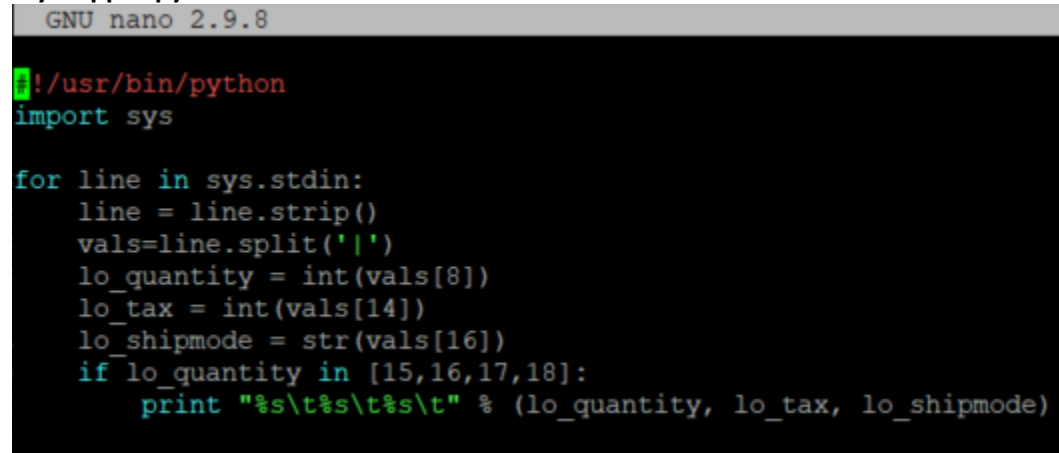
Counters:
Total records written : 11
Total bytes written : 119
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0
```



## Part 4: Hadoop Streaming

```
SELECT lo_shipmode, STDDEV(lo_tax)
FROM lineorder
WHERE lo_quantity BETWEEN 15 AND 18
GROUP BY lo_shipmode;
```

### myMapper.py

A screenshot of a terminal window showing the GNU nano 2.9.8 text editor. The editor is open to a file named myMapper.py. The script starts with a shebang line `#!/usr/bin/python` and imports the `sys` module. It then enters a `for` loop that iterates over `sys.stdin`. Inside the loop, it strips the line, splits it by the pipe character `|`, and extracts the `lo_quantity`, `lo_tax`, and `lo_shipmode` fields. It then checks if `lo_quantity` is in the range [15, 16, 17, 18]. If so, it prints the fields in a tab-separated format: `print "%s\t%s\t%s\t" % (lo_quantity, lo_tax, lo_shipmode)`.

```
GNU nano 2.9.8
#!/usr/bin/python
import sys

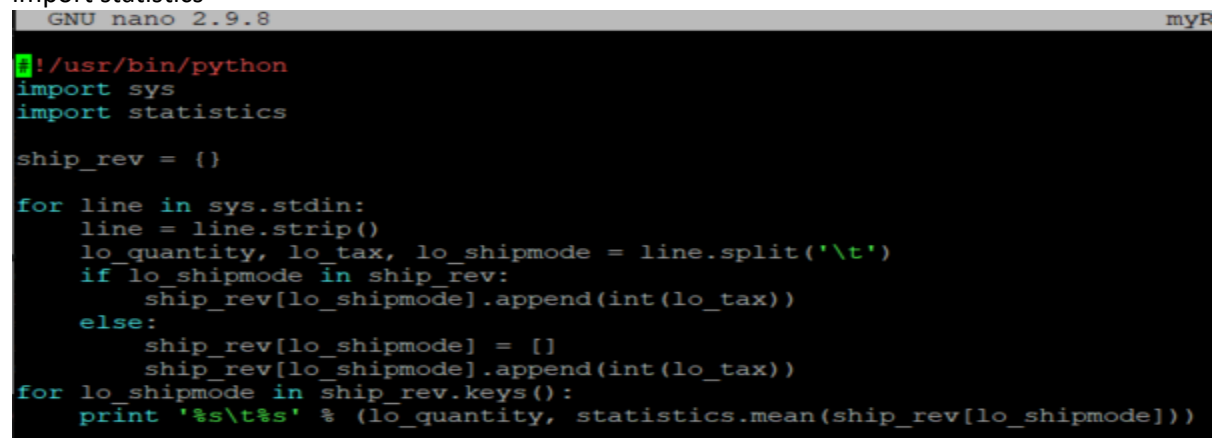
for line in sys.stdin:
    line = line.strip()
    vals=line.split('|')
    lo_quantity = int(vals[8])
    lo_tax = int(vals[14])
    lo_shipmode = str(vals[16])
    if lo_quantity in [15,16,17,18]:
        print "%s\t%s\t%s\t" % (lo_quantity, lo_tax, lo_shipmode)
```

```
#!/usr/bin/python
import sys
```

```
for line in sys.stdin:
    line = line.strip()
    vals=line.split('|')
    lo_quantity = int(vals[8])
    lo_tax = int(vals[14])
    lo_shipmode = str(vals[16])
    if lo_quantity in [15,16,17,18]:
        print "%s\t%s\t%s\t" % (lo_quantity, lo_tax, lo_shipmode)
```

### myReducer.py

```
#!/usr/bin/python
import sys
import statistics
```

A screenshot of a terminal window showing the GNU nano 2.9.8 text editor. The editor is open to a file named myReducer.py. The script starts with a shebang line `#!/usr/bin/python` and imports the `sys` and `statistics` modules. It initializes an empty dictionary `ship_rev = {}`. It then enters a `for` loop that iterates over `sys.stdin`. Inside the loop, it strips the line, splits it by the tab character `\t`, and extracts the `lo_quantity`, `lo_tax`, and `lo_shipmode` fields. It then checks if `lo_shipmode` is in the `ship_rev` dictionary. If so, it appends the `lo_tax` value to the list for that `lo_shipmode`. If not, it initializes a new list for that `lo_shipmode` and appends the `lo_tax` value. Finally, it prints the `lo_quantity` and the mean of the `lo_tax` values for each `lo_shipmode`: `print '%s\t%s' % (lo_quantity, statistics.mean(ship_rev[lo_shipmode]))`.

```
GNU nano 2.9.8 myR
#!/usr/bin/python
import sys
import statistics

ship_rev = {}

for line in sys.stdin:
    line = line.strip()
    lo_quantity, lo_tax, lo_shipmode = line.split('\t')
    if lo_shipmode in ship_rev:
        ship_rev[lo_shipmode].append(int(lo_tax))
    else:
        ship_rev[lo_shipmode] = []
        ship_rev[lo_shipmode].append(int(lo_tax))
for lo_shipmode in ship_rev.keys():
    print '%s\t%s' % (lo_quantity, statistics.mean(ship_rev[lo_shipmode]))
```

```
ship_rev = {}
```

```
for line in sys.stdin:
```

```
    line = line.strip()
```

```
    lo_quantity, lo_tax, lo_shipmode = line.split('\t')
```

```
    if lo_shipmode in ship_rev:
```

```
        ship_rev[lo_shipmode].append(int(lo_tax))
```

```
    else:
```

```
        ship_rev[lo_shipmode] = []
```

```
        ship_rev[lo_shipmode].append(int(lo_tax))
```

```
for lo_shipmode in ship_rev.keys():
```

```
    print '%s\t%s' % (lo_quantity, statistics.mean(ship_rev[lo_shipmode]))
```

## Testing

```
[ec2-user@ip-172-31-74-226 hadoop-2.6.4]$ cat lineorder.tbl | python myMapper.py | sort -n | python myReducer.py
18      4.0060694485
18      3.98971714834
18      4.01270286284
18      3.99945858148
18      3.9905022537
18      4.01043877759
18      4.00609596795
[ec2-user@ip-172-31-74-226 hadoop-2.6.4]$
```

## Output

```
hadoop jar hadoop-streaming-2.6.4.jar -input /user/ec2-user/part4 -output /data/midterm4 -mapper
myMapper.py -reducer myReducer.py -file myReducer.py -file myMapper.py
```

```
20/10/26 03:15:21 INFO mapreduce.Job: Job job_1603645671591_0024 completed successfully
20/10/26 03:15:21 INFO mapreduce.Job: Counters: 52
  File System Counters
    FILE: Number of bytes read=6382232
    FILE: Number of bytes written=13424553
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=594329915
    HDFS: Number of bytes written=117
    HDFS: Number of read operations=18
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Failed map tasks=7
    Killed map tasks=1
    Launched map tasks=12
    Launched reduce tasks=1
    Other local map tasks=7
    Data-local map tasks=5
    Total time spent by all maps in occupied slots (ms)=135617
    Total time spent by all reduces in occupied slots (ms)=11691
    Total time spent by all map tasks (ms)=135617
    Total time spent by all reduce tasks (ms)=11691
    Total vcore-milliseconds taken by all map tasks=135617
    Total vcore-milliseconds taken by all reduce tasks=11691
    Total megabyte-milliseconds taken by all map tasks=138871808
    Total megabyte-milliseconds taken by all reduce tasks=11971584
  Map-Reduce Framework
    Map input records=6001215
    Map output records=480357
    Map output bytes=5421512
    Map output materialized bytes=6382256
    Input split bytes=530
    Combine input records=0
    Combine output records=0
```



```
[ec2-user@ip-172-31-74-226 hadoop-2.6.4]$ hadoop fs -ls /data/midterm4/
Found 2 items
-rw-r--r--  2 ec2-user supergroup      0 2020-10-26 03:15 /data/midterm4/_SUCCESS
-rw-r--r--  2 ec2-user supergroup    117 2020-10-26 03:15 /data/midterm4/part-00000
[ec2-user@ip-172-31-74-226 hadoop-2.6.4]$ hadoop fs -cat /data/midterm4/part-00000
18      4.0060694485
18      3.98971714834
18      4.01270286284
18      3.99945858148
18      3.9905022537
18      4.01043877759
18      4.00609596795
[ec2-user@ip-172-31-74-226 hadoop-2.6.4]$
```